

# Primal-Dual Learning for Resource Allocation in Low Latency Wireless Systems

Mark Eisen\*   Mohammad M. Rashid†   Dave Cavalcanti†   Alejandro Ribeiro\*

**Abstract**—We consider the problem of scheduling transmissions over low-latency wireless communication links to control various control systems. Low-latency requirements are critical in developing wireless technology for industrial control, but are inherently challenging to meet while also maintaining reliable performance. An alternative to ultra reliable low latency communications (URLLC) is a framework in which reliability is adapted to control system demands. We formulate the control-aware scheduling problem as a constrained statistical optimization problem in which the optimal scheduler is a function of current control and channel states. The scheduler is parameterized with a deep neural network, and the constrained problem is solved using techniques from primal-dual learning. Furthermore, these updates are model-free in that they do not require explicit knowledge of channels models or performance metrics. The resulting control-aware deep scheduler is evaluated in empirical simulations and strong performance is shown relative to other heuristic scheduling methods.

**Index Terms**—wireless control, low-latency, deep learning, primal-dual

## I. INTRODUCTION

The recent advances in wireless technology and automation have given rise to efforts in integrating wireless communications in autonomous environments, particularly in industrial control settings where the scale of wired networks is proving increasingly costly [1], [2]. The analysis of control systems operating over wireless communication links is thus an integral part in enabling these wireless industrial automation applications. However, the performance specifications of these applications demands the design of wireless networks that can meet both the high reliability and low latency demands of the system [1], [3]–[6]. The problem of ultra reliable low latency communications (URLLC) is inherently challenging as the physical medium of wireless communication trades off these two performance criteria, making it hard to meet both reliability and latency demands.

One promising direction in enabling low latency communications involves specific developments in radio resource allocation, or scheduling. For low latency applications, traditional delay-aware schedulers [7]–[9] have been employed, in addition to more recent URLLC techniques based on various forms of diversity [5], [10]–[12]—all of which are agnostic to the control system. However, due to the physical limitations of the wireless channel, it is often necessary to use information from

the control system to make proper use of scheduling resources in meeting latency requirements. While there exist numerous ways in which control system information is incorporated into “control-aware” scheduling methods [13]–[19], these are agnostic to latency requirements of the system. More recent work [20] looks at heuristic based scheduling methods that are both control and latency aware.

Aside from traditional heuristic based scheduling methods, tools from machine learning are being incorporated into making intelligent scheduling and resource allocation decisions. The work in [21] builds a framework for solving a generic set of resource allocation problems by interpreting resource allocation as a constrained statistical learning problem. This leads a natural use of learning models, such as deep neural networks (DNNs), for designing schedulers. Recent advancements apply techniques from both reinforcement learning and deep learning for control-aware scheduling in simple systems [18], [19]. The goal in this work is to use the framework of [21] to employ so-called primal-dual deep learning to tackle develop novel scheduling techniques that are both control-aware and latency constrained.

This paper is organized as follows. We formulate the wireless control system in which state information is communicated to the control over a wireless channel as a switched dynamical system (Section II). We further discuss the scheduling architecture under which we can allocate channels and data rates between plants. With this formulation, we formulate the optimal scheduling problem that minimizes a control cost under latency constraints (Section II-A). As this problem is a statistical learning problem, it can be parameterized with a deep neural network (Section II-B). To handle the constraints, we convert the problem to the Lagrangian dual problem and present a primal-dual learning method to find the optimal parameters (Section III). We further discuss ways in which the primal-dual method can be approximated without explicit model knowledge (Section III-A). The performance of the learned control-aware scheduling method is analyzed in a numerical simulation of a control system with latency constrained and compared other heuristic scheduling methods (Section IV).

## II. WIRELESS CONTROL SYSTEMS

We consider a series of  $m$  control systems operating over a shared wireless channels. The state of plant  $i$  at control cycle index  $k$  is given by the variable  $\mathbf{x}_i^k \in \mathbb{R}^p$ . At each control/scheduling cycle, the sensor measures the state  $\mathbf{x}_i^k$  and transmits it over a wireless channel to a common base station

Supported by Intel Science and Technology Center for Wireless Autonomous Systems. The authors are with the (\*)Department of Electrical and Systems Engineering, University of Pennsylvania and (†)Wireless Communications Research, Intel Corporation. Email: maeisen@seas.upenn.edu, mamun.rashid@intel.com, dave.cavalcanti@intel.com, aribeiro@seas.upenn.edu.




Fig. 1: Scheduling architecture of  $m = 4$  users—colored in green, blue, red, and yellow—across  $n = 3$  channels.

(BS) that is co-located with the controller. Given the state information, the controller determines the necessary control input is fed back to the plant. This is referred to as the closed-loop configuration of the control cycle. Given the noisy nature of the wireless channel, there is the potential for the communications packet containing the state information to be dropped, resulting in an open-loop configuration of the control cycle. We may model the linear dynamics of the wireless control system for plant  $i$  as

$$\mathbf{x}_i^{k+1} = \begin{cases} \hat{\mathbf{A}}_i \mathbf{x}_i^k + \mathbf{w}^k & \text{if packet received} \\ \mathbf{A}_i \mathbf{x}_i^k + \mathbf{w}^k & \text{otherwise} \end{cases}, \quad (1)$$

where  $\hat{\mathbf{A}}_i \in \mathbb{R}^{p \times p}$  is the closed loop gain,  $\mathbf{A}_i \in \mathbb{R}^{p \times p}$  is the open loop gain, and  $\mathbf{w}^k \in \mathbb{R}^p$  is zero-mean i.i.d. disturbance process with covariance  $\mathbf{W}$ . The closed loop and open loop gains may reflect, e.g., controlled dynamics using accurate and estimated state information, respectively. We assume that the closed loop gains are preferable to the open loop gain, i.e.  $\lambda_{\max}(\hat{\mathbf{A}}_i) < \lambda_{\max}(\mathbf{A}_i)$ . Further note this model restricts its attention to wireless connections in uplink of the control loop, while downlink is assumed to occur over an ideal channel—i.e. no packet drops.

Given this dynamical model of the wireless control systems, the communications goal is to allocate radio resources among the various plants to maintain strong performance across all the systems. To do so, we present a generic frequency and time division multiplexing scheduling architecture with which the BS allocates scheduling resources to the plants. A scheduling window occupies the uplink of a single cycle in the control loop; the total length of this scheduling window is subject to a tight low-latency bound. Transmissions are scheduled by the BS across  $n$  different channels occupying different frequency bands. Multiple transmissions scheduled in the same channel will occur in sequence, while transmissions scheduled in different channels may occur simultaneously. Denote by  $\varsigma_i \in \{0, 1\}^n$  a binary vector whose  $j$ th element  $\varsigma_{i,j}$  is 1 if the  $i$ th device transmits in the  $j$ th channel, and 0 otherwise. Further denote for each device a data rate selection  $\mu_i \in [\mu_{\min}, \mu_{\max}]$ . These two scheduling parameters together define the scheduling decision made for the  $i$  plant. This architecture reflects generalizes many practical scheduling-based multiple access wireless protocols, such as those used in LTE [22],

5G [23], and next-generation WiFi IEEE 802.11ax [24]. An illustration of  $m = 4$  users making multiple transmission across  $n = 3$  channels is shown in Figure 1.

The achieved communications performance by a given scheduling decision can be formulated as follows. We first define  $\mathbf{h}_i^k \in \mathbb{R}_+^n$  to be the set of fading channel states experienced by device  $i$  at cycle  $k$ , where the  $j$  element  $h_{i,j}^k$  is the fading channel gain in channel  $j$ . We assume that these channel conditions do not change over the course of a scheduling window. In any given channel with fading state  $\mathbf{h}$ , we define a function  $q(\mathbf{h}, \mu)$  that returns the packet delivery rate (PDR), or the probability of successful transmission of the packet, when transmitting with data rate  $\mu$ . Likewise, we define a function  $\tau(\mu)$  that returns the transmission time to transmit a packet of fixed length with data rate  $\mu$ . These two functions play a critical role in designing low-latency wireless control systems, as they allow us to explore the trade-off between PDR and transmission time and the resulting effect on control system performance. We may consider that the functions  $q(\mathbf{h}, \mu)$  and  $\tau(\mu)$  both get smaller as we increase data rate  $\mu$ , i.e.

$$\mu' > \mu \implies q(\mathbf{h}, \mu) \leq q(\mathbf{h}, \mu'), \quad \tau(\mu') \leq \tau(\mu). \quad (2)$$

Thus, by increasing the data rate we may reduce the transmission time to satisfy latency constraints, but at the cost of control system performance, as illustrated by the switched dynamics in (1).

#### A. Optimal scheduling design

We are interested in designing scheduling policies that optimize control performance, subject to the strict low latency constraints of the system. To do so, we first formulate the global control-based performance given a scheduling decision. Collect in the matrix  $\Sigma \in \{0, 1\}^{n \times m}$  all of the channel transmission vectors  $\varsigma_i$  for  $i = 1, \dots, m$  and collect in the vector  $\boldsymbol{\mu} \in [\mu_{\min}, \mu_{\max}]^m$  the data rates  $\mu_i$  for  $i = 1, \dots, m$ . Given that a device may transmit in multiple channels within a single scheduling cycle, the probability of successful transmission can be given as the probability that the transmission was successful in at least one channel, i.e.

$$\tilde{q}(\mathbf{h}_i, \varsigma_i, \mu_i) := 1 - \prod_{j=1}^n (1 - \varsigma_{i,j} q(h_{i,j}, \mu_i)). \quad (3)$$

The total delivery rate in (3) can be viewed as the probability of receiving the packet and experiencing the closed loop dynamics in (1). Now, to evaluate the performance of a given plant at a particular state  $\mathbf{x}$ , define a quadratic Lyapunov function  $L_i(\mathbf{x}) := \mathbf{x}^T \mathbf{P}_i \mathbf{x}$  with some positive definite matrix  $\mathbf{P}_i \in \mathbb{R}^{p \times p}$ . Such a function can be used to evaluate performance or stability of the control system. Because the control system evolves in a random manner, the cost of a given scheduling decision  $\{\varsigma_i, \mu_i\}$  for the  $i$ th plant can be formulated as the *expected future Lyapunov cost* under such a

schedule. As the probability of closing the loop in (1) is given by  $\tilde{q}(\mathbf{h}_i^k, \varsigma_i, \mu_i)$ , we may write this expected future cost as

$$\begin{aligned} J_i(\mathbf{x}_i, \mathbf{h}_i, \varsigma_i, \mu_i) &:= \mathbb{E} [L_i(\mathbf{x}_i^{k+1}) \mid \mathbf{x}_i^k = \mathbf{x}_i, \mathbf{h}_i^k = \mathbf{h}_i] \quad (4) \\ &= \tilde{q}(\mathbf{h}_i, \varsigma_i, \mu_i) (\hat{\mathbf{A}}_i \mathbf{x}_i)^T \mathbf{P}_i (\hat{\mathbf{A}}_i \mathbf{x}_i) + \\ &\quad (1 - \tilde{q}(\mathbf{h}_i, \varsigma_i, \mu_i)) (\hat{\mathbf{A}}_i \xi)^T \mathbf{P}_i (\hat{\mathbf{A}}_i \xi) \\ &\quad + \text{Tr}(\mathbf{P}_i \mathbf{W}_i). \end{aligned}$$

Observe that the local control cost for the  $i$ th plant  $J_i(\mathbf{x}_i^k, \mathbf{h}_i^k, \varsigma_i, \mu_i)$  is a function of both the system *states*—the fading channel  $\mathbf{h}_i^k$  and control state  $\mathbf{x}_i^k$ —and the scheduler *actions*—channel selection  $\varsigma_i$  and data rate  $\mu_i$ . The objective is to choose the actions  $\varsigma_i$  and  $\mu_i$  that minimizes the cost relative to states  $\mathbf{h}_i^k$  and  $\mathbf{x}_i^k$ .

In addition to minimizing a control cost, we must make scheduling decisions that respect the low-latency requirements of the system. To formulate this constraint, consider the *total* time of a global scheduling decision  $\Sigma, \mu$  of channel  $j$  as the sum of all active transmissions, i.e.

$$\tilde{\tau}_j(\Sigma, \mu) := \sum_{i=1}^m \varsigma_{i,j} \tau(\mu_i). \quad (5)$$

Combining all the local costs for plants  $i = 1, \dots, m$  in (4) with the latency constraints for all channels  $j = 1, \dots, n$  in (5), we may define the optimal scheduling design problem. Because we are interested in long-term, or average, performance across random channels and control states, we optimize with respect to expected costs and probabilistic constraints. Collect all channel vectors  $\mathbf{h}_i$  in a matrix  $\mathbf{H} \in \mathbb{R}_+^{n \times m}$  and states  $\mathbf{x}_i$  in a matrix  $\mathbf{X} \in \mathbb{R}^{p \times n}$ . Consider a scheduling policy  $\mathbf{p}(\mathbf{H}, \mathbf{X}) := \{\Sigma, \mu\}$  that, given a set of channel states  $\mathbf{H}$  and control states  $\mathbf{X}$ , returns a schedule defined by the channel selection matrix  $\Sigma$  and data rate selection vector  $\mu$ . The optimal low-latency constrained scheduling policy for the wireless control systems is the one which solves the program

$$\begin{aligned} J^* &:= \min_{\mathbf{p}(\mathbf{H}, \mathbf{X})} \mathbb{E}_{\mathbf{H}, \mathbf{X}} \left[ \sum_{i=1}^m J_i(\mathbf{x}_i, \mathbf{h}_i, \varsigma_i, \mu_i) \right], \quad (6) \\ \text{s. t. } &\mathbb{P}_{\mathbf{H}, \mathbf{X}} (\tilde{\tau}_j(\Sigma, \mu) \leq t_{\max}) \geq 1 - \delta \quad \forall j, \\ &\mathbf{p}(\mathbf{H}, \mathbf{X}) := \{\Sigma \in \{0, 1\}^{n \times m}, \mu \in [\mu_{\min}, \mu_{\max}]^m\}. \end{aligned}$$

In (6), we minimize the average cost over the distribution of channel and control states, subject to the condition that the probability of violating the latency constraint over the distribution of states is less than some small value  $\delta$ . The above scheduling problem can be viewed as a constrained statistical learning problem—a connection made for a more generic class of resource allocation problems in [21]. While such a problem characterizes the optimal scheduling decision for the latency-constraint wireless control system, finding solutions to such a problem is a significant challenge. This is due to a number of complexities in (6), namely: (i) it requires functional optimization, (ii) it contains explicit constraints, and (iii) we typically do not have analytic forms for the functions and distributions in (6). The first of these complexities can be resolved using a standard technique in statistical learning, discussed next in

Section II-B. The latter two of these complexities are discussed and resolved later in Sections III and III-A, respectively.

### B. Deep learning parameterization

The scheduling problem in (6) is computationally challenging because it requires finding a policy—or *function*— $\mathbf{p}(\mathbf{H}, \mathbf{X})$ . In statistical learning, or regression, problems the regression function is replaced by some given parameterization  $\phi(\mathbf{H}, \mathbf{X}, \theta)$  that is defined with some finite dimensional parameter  $\theta \in \mathbb{R}^q$ . Their exist a wide variety of choices of this parameterization, but in modern machine learning problems the *deep neural network* (DNN) is commonly employed. This is due to the fact the DNN can be shown both empirically and analytically to contain strong representative power and generalization ability, meaning that it can approximate almost any function well. A DNN is defined as a composition of  $L$  layers, each of which consisting of a linear operation followed by a point-wise nonlinearity—also known as an activation function. More specifically, the layer  $l$  is defined by the linear operation  $\mathbf{W}_l \in \mathbb{R}^{q_{l-1} \times q_l}$  followed by a non-linear activation function  $\sigma_l : \mathbb{R}^{q_l} \rightarrow \mathbb{R}^{q_l}$ . Common choices of activation functions  $\sigma_l$  include a sigmoid function or a rectifier function (commonly referred to as ReLu). Given an input from the  $l-1$  layer  $\mathbf{w}_{l-1} \in \mathbb{R}^{q_{l-1}}$ , the resulting output  $\mathbf{w}_l \in \mathbb{R}^{q_l}$  is then computed as  $\mathbf{w}_l := \sigma_l(\mathbf{W}_l \mathbf{w}_{l-1})$ . The full DNN-parameterization of the scheduling policy is then defined as an  $L$ -layer DNN whose input at the initial layer is the concatenation of states  $\mathbf{w}_0 := [\text{vec}(\mathbf{H}); \text{vec}(\mathbf{X})]$ , i.e.

$$\phi(\mathbf{h}, \theta) := \sigma_L(\mathbf{W}_L(\sigma_{L-1}(\mathbf{W}_{L-1}(\dots(\sigma_1(\mathbf{W}_1 \mathbf{w}_0)))))). \quad (7)$$

The parameter vector  $\theta \in \mathbb{R}^q$  that defines the DNN is then the entries of  $\{\mathbf{W}_l\}_{l=1}^L$  with  $q = \sum_{l=1}^{L-1} q_l q_{l+1}$ . Further note that we can easily construct an activation function at the final layer  $\sigma_L$ —or the *output layer*—such that the outputs  $\phi(\mathbf{h}, \theta)$  are in the space  $\{0, 1\}^{n \times m} \times [\mu_{\min}, \mu_{\max}]$  that contains possible schedules. With this DNN parameterization, the control-aware scheduling problem can be rewritten as

$$\begin{aligned} J_\phi^* &:= \min_{\theta \in \mathbb{R}^q} \mathbb{E}_{\mathbf{H}, \mathbf{X}} \left[ \sum_{i=1}^m J_i(\mathbf{x}_i, \mathbf{h}_i, \alpha_i, \mu_i) \right], \quad (8) \\ \text{s. t. } &\mathbb{P}_{\mathbf{H}, \mathbf{X}} (\tilde{\tau}_j(\Sigma, \mu) \leq t_{\max}) \geq 1 - \delta \quad \forall j, \\ &\phi(\mathbf{H}, \mathbf{X}, \theta) := \{\Sigma \in \{0, 1\}^{n \times m}, \mu \in [\mu_{\min}, \mu_{\max}]^m\}. \end{aligned}$$

Observe in (8) that the optimization is performed over  $\theta$  rather than the scheduling policy directly. In other words, we look for the interlayer weights that define a DNN that minimizes the total control cost while satisfying the latency constraints. We proceed then to discuss a learning method that can find solutions to the constrained optimization problem in (8).

## III. PRIMAL-DUAL LEARNING

Finding the DNN layer weights  $\theta$  that provide good solutions to (8) requires the solving of a constraint learning problem. The standard approach of gradient-based optimization methods cannot be applied directly here to the presence of the latency constraints. To proceed then, we must formulate an

unconstrained problem that captures the form of (8). A naive penalty-based reformulation will introduce a similar but fundamentally different problem, so we thus opt for constructing a Lagrangian dual problem. For notational convenience, moving forward we employ the following shorthands for the state variables, aggregate Lyapunov function, latency constraint functions, respectively:

$$\mathbf{w} := [\text{vec}(\mathbf{H}); \text{vec}(\mathbf{X})], \quad (9)$$

$$f(\phi(\mathbf{w}, \boldsymbol{\theta}), \mathbf{w}) := \sum_{i=1}^m J_i(\mathbf{x}_i, \mathbf{h}_i, \boldsymbol{\varsigma}_i, \boldsymbol{\mu}_i), \quad (10)$$

$$g_j(\phi(\mathbf{w}, \boldsymbol{\theta}), \mathbf{w}) := \mathbb{I}[\tilde{\tau}_j(\boldsymbol{\Sigma}, \boldsymbol{\mu}) \leq t_{\max}] - (1 - \delta) \quad (11)$$

We introduce the nonnegative dual variables  $\boldsymbol{\lambda} \in \mathbb{R}_+^n$  associated with the vector of constraint functions  $\mathbf{g}(\mathbf{p}(\mathbf{w}, \boldsymbol{\theta}), \mathbf{w}) := [g_1(\cdot); \dots; g_n(\cdot)]$ , and form the Lagrangian as

$$\mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\lambda}) := \mathbb{E}_{\mathbf{w}} [f(\phi(\mathbf{w}, \boldsymbol{\theta}), \mathbf{w}) - \boldsymbol{\lambda}^T \mathbf{g}(\phi(\mathbf{w}, \boldsymbol{\theta}), \mathbf{w})]. \quad (12)$$

The Lagrangian in (12) penalizes constraint violation through the second term. Note, however, that the penalty is scaled by the dual parameter  $\boldsymbol{\lambda}$ . The so-called Lagrangian dual problem is one in which both the primal variable  $\boldsymbol{\theta}$  is simultaneously minimized while the dual parameter  $\boldsymbol{\lambda}$  is maximized. Such a problem can be written with the saddle point formulation

$$D_{\phi}^* := \max_{\boldsymbol{\lambda} \geq 0} \min_{\boldsymbol{\theta}} \mathcal{L}(\boldsymbol{\theta}, \boldsymbol{\lambda}). \quad (13)$$

The dual optimum  $D_{\phi}^*$  is the best approximation of the form in (12) we can have of  $J_{\phi}^*$ . In fact, under some standard assumptions on the problem and assuming a sufficiently dense DNN architecture, we can formally bound the difference between  $D_{\phi}^*$  and  $J^*$  to be proportional to the approximation capacity of the DNN  $\phi(\mathbf{H}, \mathbf{X}, \boldsymbol{\theta})$ —see [21] for details on this result. Thus, we may say that, up to some approximation, solving the unconstrained problem in (13) is equivalent to solving the constrained problem in (8).

With the unconstrained saddle point problem in (13), we may perform standard gradient-based optimization methods to obtain solutions. The max-min structure necessitates the use of a *primal-dual* learning method, in which we iteratively update both the primal and dual variable in (12) to find a local stationary point of the KKT conditions of (8). Consider a learning iteration index  $t = 0, 1, \dots$  over which we define a sequence of primal variables  $\{\boldsymbol{\theta}_t\}$  and dual variables  $\{\boldsymbol{\lambda}_t\}$ . At index  $t$ , we determine the value of next primal iterate  $\mathbf{x}_{t+1}$  by adding to the current iterates the corresponding partial gradients of the Lagrangian in (12)  $\nabla_{\boldsymbol{\theta}} \mathcal{L}$ , i.e.,

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \alpha_t \nabla_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{w}} [f(\phi(\mathbf{w}, \boldsymbol{\theta}_t), \mathbf{w}) - \boldsymbol{\lambda}_t^T \mathbf{g}(\phi(\mathbf{w}, \boldsymbol{\theta}_t), \mathbf{w})], \quad (14)$$

where we introduce  $\alpha_t > 0$  as a scalar step size. We subsequently perform a corresponding partial gradient update to compute the dual iterate  $\boldsymbol{\lambda}_{t+1}$ , i.e.

$$\boldsymbol{\lambda}_{t+1} = [\boldsymbol{\lambda}_t - \beta_t \mathbb{E}_{\mathbf{w}} \mathbf{g}(\phi(\mathbf{w}, \boldsymbol{\theta}_{t+1}), \mathbf{w})]_+, \quad (15)$$

with associated step size  $\beta_t > 0$ . Observe in (15) that we additionally project onto the positive orthant to maintain the nonnegative constraint on  $\boldsymbol{\lambda}$ . The gradient primal-dual updates in (14) and (15) successively move the primal and dual variables towards maximum and minimum points of the Lagrangian function, respectively.

#### A. Model-free updates

The updates in (14)-(15) cannot, in general, be applied exactly. To see this, observe that computing the gradients in (14) requires computing the gradient of  $J_i(\cdot)$ —which depends on PDR function  $\tilde{q}(\cdot)$ —and the gradient of an indicator of transmission length function  $\tilde{\tau}(\cdot)$ . In practical systems, we do not typically have easily available analytic forms for these functions to take gradients. Furthermore, both the updates in (14) and (15) require to take the expectation over the distribution of states  $\mathbf{x}$  and  $\mathbf{h}$ . These, too, are often unknown in practice. However, there exist standard ways of approximating the updates with stochastic, *model-free* updates that do not require such knowledge. Most popular among these is the policy gradient approximation [25].

To compute a policy gradient update, we consider the scheduling parameters  $\boldsymbol{\Sigma}$  and  $\boldsymbol{\mu}$  are drawn stochastically from a distribution with given form  $\pi_{\phi(\mathbf{w}, \boldsymbol{\theta})}$  whose parameters are given by the output of the DNN  $\phi(\mathbf{w}, \boldsymbol{\theta})$ —e.g. the mean and variance of a normal distribution. Using such a stochastic policy, it can be shown that an unbiased estimators of the gradients in (14) and (15) can be formed as,

$$\widehat{\nabla}_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{w}} f(\phi(\mathbf{w}, \boldsymbol{\theta}), \mathbf{w}) = f(\hat{\mathbf{p}}_{\boldsymbol{\theta}}, \hat{\mathbf{w}}) \nabla_{\boldsymbol{\theta}} \log \pi_{\phi(\hat{\mathbf{w}}, \boldsymbol{\theta})}(\hat{\mathbf{p}}_{\boldsymbol{\theta}}) \quad (16)$$

$$\widehat{\nabla}_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{w}} \mathbf{g}(\phi(\mathbf{w}, \boldsymbol{\theta}), \mathbf{w}) = \mathbf{g}(\hat{\mathbf{p}}_{\boldsymbol{\theta}}, \hat{\mathbf{w}}) \nabla_{\boldsymbol{\theta}} \log \pi_{\phi(\hat{\mathbf{w}}, \boldsymbol{\theta})}(\hat{\mathbf{p}}_{\boldsymbol{\theta}})^T \quad (17)$$

$$\widehat{\mathbb{E}}_{\mathbf{w}} \mathbf{g}(\phi(\mathbf{w}, \boldsymbol{\theta}), \mathbf{w}) = \mathbf{g}(\hat{\mathbf{p}}_{\boldsymbol{\theta}}, \hat{\mathbf{w}}), \quad (18)$$

where  $\hat{\mathbf{w}}$  is a sampled state and  $\hat{\mathbf{p}}_{\boldsymbol{\theta}}$  is a sample drawn from the distribution  $\pi_{\phi(\hat{\mathbf{w}}, \boldsymbol{\theta})}$ . In practice, we may reduce the variance of these unbiased estimates by taking  $B$  samples and averaging. Note that the updates here only require taking gradients of the log likelihoods rather than of the functions themselves. Thus, we can replace the updates in (14) and (15) with their model free counterparts, i.e.

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \alpha_t \widehat{\nabla}_{\boldsymbol{\theta}} \mathbb{E}_{\mathbf{w}} [f(\phi(\mathbf{w}, \boldsymbol{\theta}_t), \mathbf{w}) - \boldsymbol{\lambda}_t^T \mathbf{g}(\phi(\mathbf{w}, \boldsymbol{\theta}_t), \mathbf{w})], \quad (19)$$

$$\boldsymbol{\lambda}_{t+1} = [\boldsymbol{\lambda}_t - \beta_t \widehat{\mathbb{E}}_{\mathbf{w}} \mathbf{g}(\phi(\mathbf{w}, \boldsymbol{\theta}_{t+1}), \mathbf{w})]_+. \quad (20)$$

The complete primal-dual learning algorithm is summarized in Algorithm 1. We conclude with a brief remark on state sampling.

**Remark 1:** In the gradient estimations in, e.g. (16), we sample both the control states  $\mathbf{x}$  and channel states  $\mathbf{h}$ . This assumes that such samples can be drawn i.i.d. While this may generally be true for the channel states  $\mathbf{h}$ , it will not be generally be true for the control states  $\mathbf{x}$  in practice, due to the fact that the states evolve based on the switched dynamics in (1), which itself depends on the scheduling actions taken. A more precise way to model the statistics of the control

---

**Algorithm 1** Model-Free Primal-Dual Learning

---

1: **Parameters:** Policy model  $\phi(\mathbf{h}, \theta)$  and distribution form  $\pi_{\mathbf{h}, \theta}$   
2: **Input:** Initial states  $\theta_0, \lambda_0$   
3: **for**  $t = 0, 1, 2, \dots$  **do** {main loop}  
4:   Draw samples  $\{\hat{\theta}, \hat{\mathbf{h}}\}$ , or in batches of size  $B$   
5:   Compute policy gradients [ c.f. (16)-(18)]  
6:   Update primal and dual variables  
$$\theta_{t+1} = \theta_t - \alpha_t \widehat{\nabla_{\theta}} \mathbb{E}_{\mathbf{w}} [f(\phi(\mathbf{w}, \theta_t), \mathbf{w}) - \lambda_t^T \mathbf{g}(\phi(\mathbf{w}, \theta_t), \mathbf{w})], [cf.(19)]$$
$$\lambda_{t+1} = [\lambda_t - \beta_t \widehat{\mathbb{E}_{\mathbf{w}}} \mathbf{g}(\phi(\mathbf{w}, \theta_{t+1}), \mathbf{w})]_+ [cf.(20)]$$
  
7: **end for**

---

states would be with a Markov decision process (MDP). The generalization of the presented techniques for this setting make up what is known as *reinforcement learning* algorithms. In this work, we nonetheless assume that  $\mathbf{x}$  can also be drawn i.i.d. from an approximate distribution and leave the full MDP formulation as the study of future work.

#### IV. SIMULATION RESULTS

We perform a series of simulations on latency-constrained wireless control systems to evaluate the performance the learning method in and the resulting control-aware scheduling policies. We generate a series  $m = 20$  plants with closed-loop gains  $\hat{\mathbf{A}}_i \sim \text{Uniform}(0.8, 0.95)$  and open-loop gains  $\hat{\mathbf{A}}_i \sim \text{Uniform}(1.01, 1.3)$ . The variance for all system noise  $\mathbf{w}_i$  is set to be  $W = 1$ . All such plants send their state information over a shared wireless channel with  $n = 5$  independent channels with a total latency constraint of  $t_{\max} = 0.5$  ms. A latency bound of this order is typical of industrial control systems such as printing machines and presses [26]. We further assume that the states of the plants are confined to the box  $[-10, 10]$ . The DNN is given an architecture of 4 layers of sizes 256, 128, 64, and 32, all using a ReLU activation function.

With the scheduling architecture given in Figure 1 for 5 channels and 20 plants, at each control scheduling interval each plant is given a data rate  $\mu_i$  and a set of channels to transmit on. In our simulations, we use the modulation and coding schemes (MCS) of the next-generation IEEE 802.11ax Wi-Fi protocol as a representative architecture for data rate selection and packet error rate computation. As such, the continuous data rates  $\mu_i$  are selected in an interval of  $[1.6, 8]$  and rounded down to the nearest discrete MCS selection given in 802.11ax—see [24] for details on the MCS tables given in this protocol. The corresponding transmission time  $\tau(\mu)$  is then calculated assuming a fixed packet size of 100 bytes and the packet delivery rate  $q(h, \mu)$  is computed using the associated AWGN error curve (scaled by the effective SNR given channel conditions).

In Figure 2 we show the training process for the control-aware scheduler using the primal-dual scheduler given in Algorithm 1. In the left figure, we show the aggregate Lyapunov

control cost over the course of 40,000 learning iterations and in the right figure, show the transmission time of each channel (shown in different colored lines). As can be seen the in the figure, the policies initially output infeasible schedules that are beyond the latency limit but eventually converge to scheduling decisions that respect latency requirements. The control objective continues to decrease after feasibility is obtained and eventually converges to a local minimum value.

We proceed to simulate the performance of the control systems using the scheduling policy obtained at the end of the learning process and compare against a standard, round-robin scheduling policy—i.e. control-agnostic. Because the constraint in (6) is satisfied only with probability  $1 - \delta$ , we ensure latency requirements are met at runtime by restricting all transmissions that are scheduled after the latency bound. In Figure 3 we show the state values over 1000 control cycles for the DNN-based scheduler (top figure) and the round robin scheduler (bottom figure). It can be observed that, overall, the plants are kept in better states using the DNN scheduler than with the round robin scheduler. This representative simulation highlights the ability of control-aware scheduling to better meet strict latency demands and, furthermore, highlight the ability of a DNN to properly encode such a constraints and objective in its dense architecture.

#### V. CONCLUSION

In this paper we develop a control-communication control-aware approach towards scheduling for low-latency wireless control systems. To handle the challenge of achieving high reliability performance with limited scheduling resources, we formulate a control-aware scheduling problem in which reliability is adapted to control and channel states and plant dynamics. This problem takes the form of a constrained statistical learning problem, in which solutions can be found by parameterized the scheduling policy with a deep neural network and applying a primal-dual learning framework to obtain the optimal weights of the neural network. Numerical simulations showcase the ability of the primal-dual method to find scheduling policies that outperform simple control-agnostic scheduling procedures.

#### REFERENCES

- [1] Pouria Zand, Supriyo Chatterjea, Kallol Das, and Paul Havinga, “Wireless industrial monitoring and control networks: The journey so far and the road ahead,” *Journal of sensor and actuator networks*, vol. 1, no. 2, pp. 123–152, 2012.
- [2] Martin Wollschlaeger, Thilo Sauter, and Juergen Jasperneite, “The future of industrial communication: Automation networks in the era of the internet of things and industry 4.0,” *IEEE Industrial Electronics Magazine*, vol. 11, no. 1, pp. 17–27, 2017.
- [3] Anitha Varghese and Deepaknath Tandur, “Wireless requirements and challenges in industry 4.0,” in *Contemporary Computing and Informatics (IC3I)*, 2014 *International Conference on*. IEEE, 2014, pp. 634–638.
- [4] Matthew Weiner, Milos Jorgovanovic, Anant Sahai, and Borivoje Nikolić, “Design of a low-latency, high-reliability wireless communication system for control applications,” in *Communications (ICC)*, 2014 *IEEE International Conference on*. IEEE, 2014, pp. 3829–3835.




Fig. 2: Convergence of (left) objective function value, and (right) transmission time for a low-latency, control aware scheduling policy over the learning process. The DNN parameterized scheduling policy obtains feasible latency-contained schedules ( $t_{\max} = 5 \times 10^{-4}$ ) that converges to a local minimum.




Fig. 3: Evolution of plant states over 1000 control cycles using both the (top) DNN-based control aware scheduler and (b) a round robin scheduler.

[5] Petar Popovski, Jimmy J Nielsen, Cedimir Stefanovic, Elisabeth de Carvalho, Erik Strom, Kasper F Trillingsgaard, Alexandru-Sabin Bana, Dong Min Kim, Radoslaw Kotaba, Jihong Park, et al., “Wireless access for ultra-reliable low-latency communication: Principles and building blocks,” *IEEE Network*, vol. 32, no. 2, pp. 16–23, 2018.

[6] Mehdi Bennis, Mérouane Debbah, and H Vincent Poor, “Ultra-reliable and low-latency wireless communication: Tail, risk and scale,” *arXiv preprint arXiv:1801.01270*, 2018.

[7] Chengjie Wu, Mo Sha, Dolvara Gunatilaka, Abusayeed Saifullah, Chenyang Lu, and Yixin Chen, “Analysis of edf scheduling for wireless sensor-actuator networks,” in *Quality of Service (IWQoS), 2014 IEEE 22nd International Symposium of*, IEEE, 2014, pp. 31–40.

[8] Songwu Lu, Vaduvur Bhargavan, and Rayadurgam Srikant, “Fair scheduling in wireless packet networks,” *IEEE/ACM Transactions on networking*, vol. 7, no. 4, pp. 473–489, 1999.

[9] Matthew Andrews, Krishnan Kumaran, Kavita Ramanan, Alexander Stolyar, Phil Whiting, and Rajiv Vijayakumar, “Providing quality of service over a shared wireless link,” *IEEE Communications magazine*, vol. 39, no. 2, pp. 150–154, 2001.

[10] Vasuki Narasimha Swamy, Sahaana Suri, Paul Rigge, Matthew Weiner, Gireeja Ranade, Anant Sahai, and Borivoje Nikolić, “Cooperative communication for high-reliability low-latency wireless control,” in *Communications (ICC), 2015 IEEE International Conference on*, IEEE, 2015, pp. 4380–4386.

[11] Jimmy Jessen Nielsen, Rongkuan Liu, and Petar Popovski, “Ultra-reliable low latency communication using interface diversity,” *IEEE Transactions on Communications*, vol. 66, no. 3, pp. 1322–1334, 2018.

[12] Muhammad Ikram Ashraf, Chen-Feng Liu, Mehdi Bennis, Walid Saad, and Choong Seon Hong, “Dynamic resource allocation for optimized

latency and reliability in vehicular networks,” *IEEE Access*, vol. 6, pp. 63843–63858, 2018.

[13] Anton Cervin and Toivo Henningsson, “Scheduling of event-triggered controllers on a shared network,” in *Proc. of the 47th IEEE Conf. on Dec. and Control (CDC)*, 2008, pp. 3601–3606.

[14] Mohammad H Mamduhi, Domagoj Tolić, Adam Molin, and Sandra Hirche, “Event-triggered scheduling for stochastic multi-loop networked control systems with packet dropouts,” in *Decision and Control (CDC), 2014 IEEE 53rd Annual Conference on*, IEEE, 2014, pp. 2776–2782.

[15] Ling Shi, Peng Cheng, and Jiming Chen, “Optimal periodic sensor scheduling with limited resources,” *IEEE Transactions on Automatic Control*, vol. 56, no. 9, pp. 2190–2195, 2011.

[16] Duo Han, Junfeng Wu, Huanshui Zhang, and Ling Shi, “Optimal sensor scheduling for multiple linear dynamical systems,” *Automatica*, vol. 75, pp. 260–270, 2017.

[17] Konstantinos Gatsis, Miroslav Pajic, Alejandro Ribeiro, and George J. Pappas, “Opportunistic control over shared wireless channels,” *IEEE Transactions on Automatic Control*, vol. 60, no. 12, pp. 3140–3155, December 2015.

[18] Burak Demirel, Arunselvan Ramaswamy, Daniel E Quevedo, and Holger Karl, “Deepcas: A deep reinforcement learning algorithm for control-aware scheduling,” *IEEE Control Systems Letters*, vol. 2, no. 4, pp. 737–742, 2018.

[19] Alex S Leong, Arunselvan Ramaswamy, Daniel E Quevedo, Holger Karl, and Ling Shi, “Deep reinforcement learning for wireless sensor scheduling in cyber-physical systems,” *arXiv preprint arXiv:1809.05149*, 2018.

[20] Mark Eisen, Mohammad M Rashid, Konstantinos Gatsis, Dave Cavalcanti, Nageen Himayat, and Alejandro Ribeiro, “Control aware radio resource allocation in low latency wireless control systems,” *arXiv preprint arXiv:1811.00409*, 2018.

[21] Mark Eisen, Clark Zhang, Luiz FO Chamon, Daniel D Lee, and Alejandro Ribeiro, “Learning optimal resource allocations in wireless systems,” *arXiv preprint arXiv:1807.08088*, 2018.

[22] Stefania Sesia, Matthew Baker, and Issam Toufik, *LTE-the UMTS long term evolution: from theory to practice*, John Wiley & Sons, 2011.

[23] Mamta Agiwal, Abhishek Roy, and Navrati Saxena, “Next generation 5g wireless networks: A comprehensive survey,” *IEEE Communications Surveys & Tutorials*, vol. 18, no. 3, pp. 1617–1655, 2016.

[24] Jianhan Liu, R Porat, N Jindal, et al., “Ieee 802.11 ax channel model document,” *Wireless LANs, Rep. IEEE 802.11-14/0882r3*, 2014.

[25] Richard S Sutton, David A McAllester, Satinder P Singh, and Yishay Mansour, “Policy gradient methods for reinforcement learning with function approximation,” in *Advances in neural information processing systems*, 2000, pp. 1057–1063.

[26] Shehzad A Ashraf, Ismet Aktas, Erik Eriksson, Ke Wang Helmersson, and Junaid Ansari, “Ultra-reliable and low-latency communication for wireless factory automation: From LTE to 5G,” in *Emerging Technologies and Factory Automation (ETFA), 2016 IEEE 21st International Conference on*, IEEE, 2016, pp. 1–8.